

# Filipino Hate Speech Detection — Revised Technical Documentation

---

Mirrored from "ann docu.pdf" with corrections + system flow

Project Team

2025-12-13

- [1.1. Problem Statement](#)
  - [1.1.1 Identified Problem](#)
  - [1.2.1 Real-World Impact](#)
  - [1.3.1 Project Objective](#)
- [2.2. Solution Approach](#)
  - [2.1.2 Deep Neural Network Selection](#)
  - [2.2.2 Rationale for Choosing BiLSTM](#)
  - [2.3.2 Rationale for Choosing ELECTRA Transformer](#)
  - [2.4.2 System Flow \(End-to-End\)](#)
- [3.3. Dataset Information](#)
  - [3.1.3.1 Dataset Sources](#)
  - [3.2.3.2 Dataset Statistics](#)
  - [3.3.3.3 Data Validation](#)
- [4.4. Neural Network Architecture](#)
  - [4.1.4.1 BiLSTM Model Architecture](#)
  - [4.2.4.2 ELECTRA Transformer Architecture](#)
- [5.5. Training Configuration](#)
  - [5.1.5.1 Hardware Specifications](#)
  - [5.2.5.2 Software Environment](#)
  - [5.3.5.3 Training Parameters - BiLSTM](#)
  - [5.4.5.4 Training Parameters - ELECTRA](#)
- [6.6. Hyperparameter Tuning Experiments](#)
  - [6.1.6.1 Overview](#)
  - [6.2.6.2 Experiment 1: BiLSTM Baseline Configuration](#)
  - [6.3.6.3 Experiment 2: ELECTRA Transformer Fine-tuning](#)
  - [6.4.6.4 Tuning Tips](#)
- [7.7. Results and Performance](#)
  - [7.1.7.1 Model Comparison \(qualitative\)](#)
  - [7.2.7.2 Target Achievement](#)
  - [7.3.7.3 Detailed Performance Analysis - BiLSTM](#)
  - [7.4.7.4 Model Strengths and Limitations](#)
- [8.8. Tools and Technologies](#)
  - [8.1.8.1 Deep Learning Framework](#)
  - [8.2.8.2 Data and Text Processing](#)
  - [8.3.8.3 Model Training Modules](#)
  - [8.4.8.4 Evaluation and Visualization](#)
  - [8.5.8.5 Development Environment](#)

- [8.6 8.6 GUI Applications](#)
- [9.9. Model Deployment](#)
  - [9.1 9.1 Inference Pipeline](#)
  - [9.2 9.2 Application Interfaces](#)
  - [9.3 9.3 Deployment Files](#)
- [10 10. Conclusion](#)
  - [10.1 10.1 Project Achievements](#)
  - [10.2 10.2 Key Findings](#)
  - [10.3 10.3 Future Improvements](#)
  - [10.4 10.4 Lessons Learned](#)
- [11 11. Practical Notes](#)
- [12 12. Reproducibility](#)

# Filipino Hate Speech Detection

Revised Technical Documentation — Corrected Details & System Flow

Date: 2025-12-13

## 1 1. Problem Statement

---

### 1.1 1.1 Identified Problem

The proliferation of hate speech and cyberbullying on Filipino social media platforms poses a significant threat to online community safety and user well-being. The unique linguistic characteristics of Filipino online content present several challenges: Key Challenges:

- Bilingual Code-Switching: Filipino users frequently alternate between Filipino (Tagalog) and English within the same sentence.
- Colloquial Language: Heavy use of slang, abbreviations, and informal expressions.
- Cultural Context: Hate speech often requires deep cultural understanding to identify correctly.
- Limited Resources: Fewer pre-trained models and annotated datasets are available for Filipino language processing compared to English.
- Social Media Dynamics: Unique features like hashtags, mentions, and links require special handling.

### 1.2 1.2 Real-World Impact

- Psychological harm and trauma to victims.
- Toxic online environments that suppress free expression.
- Potential escalation to real-world conflicts.
- Reduced user engagement and platform trust.
- Legal and ethical concerns for platform operators.

### 1.3 1.3 Project Objective

Develop an automated hate speech detection system capable of:

- Accurately identifying hate speech

in Filipino and bilingual (Filipino-English) text. - Processing social media content with informal language patterns. - Achieving at least 50-60% accuracy on test data. - Providing deployable solutions through GUI applications.

## 2. Solution Approach

---

We mirror the original document's structure and wording while correcting technical details to match the current repository.

### 2.1 Deep Neural Network Selection

This project implements two complementary approaches for hate speech detection: Approach 1: Bidirectional LSTM (BiLSTM) - Role: Baseline model (custom implementation for Filipino hate speech). - Description: A Recurrent Neural Network optimized for sequential data. Approach 2: ELECTRA Transformer (Fine-tuned) - Role: Primary approach in this repository (transfer learning from a pre-trained Filipino model). - Description: Used as a strong benchmark and practical production path.

### 2.2 Rationale for Choosing BiLSTM

- Sequential Nature of Language: Captures temporal dependencies in sequences.
- Bidirectional Context: Understands context from past and future words simultaneously.
- Long-term Dependencies: LSTM gates retain important information and forget irrelevant details.
- Proven Performance: Strong for text classification with limited resources.
- Resource Efficiency: Fewer computational requirements vs. transformers.
- Customization: Full control to adapt to Filipino linguistic patterns.

### 2.3 Rationale for Choosing ELECTRA Transformer

- Transfer Learning: Leverages pre-training on a large Filipino corpus.
- Superior Accuracy: Demonstrated ~86% test accuracy on Filipino hate speech benchmarks.
- Contextual Understanding: Attention mechanisms capture complex relationships.
- Pre-trained Knowledge: Strong understanding of Filipino structure and semantics.

### 2.4 System Flow (End-to-End)

End-to-end ELECTRA pipeline used in this repository:

1. Data ingest (CSV or HuggingFace) → consolidate splits
2. Preprocess and add social tokens ([LINK], [MENTION], [HASHTAG] when applicable)
3. Tokenize via WordPiece and cache TensorDatasets (MD5-keyed)
4. Initialize pretrained model and tokenizer
5. Fine-tune with validation; track best checkpoint
6. Evaluate on test; export weights
7. Deploy for inference (CLI, scripts, or service)

ASCII flow:

```
[CSV / HF Datasets]
```

```
[Preprocess + Add Tokens]
|
[Tokenize (WordPiece)]
|
[Cache Datasets]
|
[Fine-Tune ELECTRA + Eval]
|
[Best Checkpoint Saved]
|
[Test + Report]
```

## 3 3. Dataset Information

---

### 3.1 3.1 Dataset Sources

Primary datasets in this repository: 1. Local Hate Speech Splits - Directory: `hatespeech/` (`train.csv`, `valid.csv`, `test.csv`) - Content: Filipino social media hate speech (binary labels `text`, `label`). 2. Filipino TikTok Hate Speech Dataset - Directory: `filipino-tiktok-hatespeech-main/data` (`train.csv`, `valid.csv`, `test.csv`) - Platform: TikTok comments; modern social media interactions. 3. Cyberbullying Tweets Dataset - File: `cyberbullying_tweets.csv` - Platform: Twitter/X; labeled cyberbullying instances. 4. Aggression Dataset - File: `aggression_parsed_dataset.csv` - Content: Aggression-related text for auxiliary experiments. 5. Optional HuggingFace Dataset - Source: `mteb/FilipinoHateSpeechClassification` via `datasetImportLib.py` - Note: Requires `huggingface-cli login`.

Note: Earlier versions referenced unified CSVs (e.g., `unified_filipino_hatespeech.csv`). In this repository, we use the datasets above.

### 3.2 3.2 Dataset Statistics

Split Ratio: - Training (~70%): Model learning. - Validation (~15%): Hyperparameter tuning. - Test (~15%): Final evaluation.

Label Distribution (observed trend): - Class 0 (Non-Hate Speech): Majority class. - Class 1 (Hate Speech): Minority class.

### 3.3 3.3 Data Validation

Bias Assessment Conducted: - Geographic Bias: Diverse Filipino regions and dialects; multiple sources reduce single-source bias. - Platform Diversity: TikTok, Twitter/X, and others for varied styles. - Temporal Bias: Data across time captures evolving slang; periodic retraining recommended. - Class Imbalance: Addressed via appropriate metrics (Precision, Recall, F1) and threshold/weighting.

Privacy Considerations: - User Anonymization: Personal identifiers removed; only text retained. - Public Content Only: Sourced from public posts. - Sensitive Information: No sensitive PII; URLs/mentions/hashtags handled via placeholders.

## 4 4. Neural Network Architecture

---

## 4.1 4.1 BiLSTM Model Architecture

Model Class: `BiLSTMHateSpeechClassifier` (baseline implementation)

Structure: 1. Input Text: Tokenization & Padding. 2. Embedding Layer: `vocab_size = 10,000`; `dim = 128`. (Corrected; not 20,000.) 3. Dropout: `rate = 0.3`. 4. Bidirectional LSTM Layer 1: `hidden = 128` per direction → output 256. 5. Bidirectional LSTM Layer 2: `hidden = 128` per direction → output 256. 6. Dropout: `rate = 0.3`. 7. Fully Connected Layer 1: 256 → 64 (ReLU). 8. Dropout: `rate = 0.3`. 9. Fully Connected Layer 2: 64 → 1 (Sigmoid at inference). 10. Output: Hate Speech Probability [0.0 - 1.0].

Layer Details: - Embedding Layer: 10,000 words × 128 dims. - LSTMs: Two stacked BiLSTMs with 128 hidden units per direction. - Dropout: 0.3 at embedding/LSTM junctions and before FC. - Fully Connected: 256→64→1. - Loss Function: `BCEWithLogitsLoss` (supports `pos_weight`). - Max Sequence Length: 128 tokens. (Corrected; not 100.)

Model Parameters (approximate): - Total: ~2.0M parameters (lower than earlier ~2.9M due to 10k vocab).

## 4.2 4.2 ELECTRA Transformer Architecture

- Model: [jcblaise/electra-tagalog-small-cased-discriminator](#)
- Structure: Encoder with multi-head self-attention; classification head for fine-tuning.
- Parameters: ~14M (small).
- Tokens: Utilizes `[CLS]`, `[SEP]`; can add `[LINK]`, `[MENTION]`, `[HASHTAG]`.
- Max Sequence Length: 128 (configurable with `--msl`).

# 5 5. Training Configuration

---

## 5.1 5.1 Hardware Specifications

- GPU: CUDA-capable GPU recommended.
- CPU: Modern multi-core processor.
- Resource Requirements (indicative): BiLSTM ~2–4 GB; ELECTRA ~8–16 GB.

## 5.2 5.2 Software Environment

- Frameworks: PyTorch 1.x/2.x (CUDA support), HuggingFace Transformers 4.x.
- Language: Python 3.8+ recommended.
- Libraries: pandas, numpy, scikit-learn, tqdm, optional Weights & Biases.

## 5.3 5.3 Training Parameters - BiLSTM

- Batch Size: 64
- Learning Rate: 0.001
- Optimizer: Adam ( $\beta_1=0.9$ ,  $\beta_2=0.999$ )
- Weight Decay: 0.0
- Epochs: 10
- Loss Function: `BCEWithLogitsLoss` (recommend `pos_weight` for imbalance)
- Threshold: tuned; 0.75 reduced false positives in analysis.
- Max Sequence Length: 128 (corrected).

## 5.4 5.4 Training Parameters - ELECTRA

- Batch Size: 32 (typical; depends on GPU)
- Learning Rate: 0.0002
- Optimizer: AdamW/Adam
- Epochs: 3
- Warmup Percentage: 0.1
- Scheduler: Linear with warmup

Trainer CLI (this repository):

```
python Filipino-Text-Benchmarks-master/train.py \
--pretrained jcblaise/electra-tagalog-small-cased-discriminator \
--train_data hatespeech/train.csv \
--valid_data hatespeech/valid.csv \
--test_data hatespeech/test.csv \
--checkpoint finetuned_model \
--msl 128 --batch_size 32 --learning_rate 2e-4 --epochs 3 \
--add_token [LINK],[MENTION],[HASHTAG]
```

Notes: Use `--data_pct` for low-resource simulation; `--label_columns` for multi-label; `--text_columns s1,s2` for sentence pairs. Caching is automatic and keyed by data paths + hyperparameters.

# 6 6. Hyperparameter Tuning Experiments

---

## 6.1 6.1 Overview

Two main architectures were considered to compare performance and resource needs.

## 6.2 6.2 Experiment 1: BiLSTM Baseline Configuration

Experiment ID: BiLSTM-001 Status: Completed - Baseline Training Configuration: - Embedding Dimension: 128 - Hidden Dimension: 128 - Layers: 2 - Dropout: 0.3 - Batch Size: 64 - Learning Rate: 0.001

Observations: - Strengths: Efficient training; good sequential modeling. - Optimizer Choice: Adam worked well without manual scheduling.

## 6.3 6.3 Experiment 2: ELECTRA Transformer Fine-tuning

Experiment ID: ELECTRA-001 Status: Configured/Run based on benchmarks Training Configuration: - Base Model: jcblaise/electra-tagalog-small-cased-discriminator - Learning Rate: 0.0002 - Epochs: 3

Performance Analysis: - Strengths: Superior accuracy (~86%) due to Filipino pre-training. - Trade-offs: Higher computational requirements.

## 6.4 6.4 Tuning Tips

- ELECTRA knobs: learning rate 2e-5–5e-4; batch 16–64; msl 128–256; warmup.
- BiLSTM knobs: embedding size, hidden size, layers, dropout; threshold, `pos_weight`.
- Optional W&B sweeps:

```
wandb sweep -p PROJECT_NAME Filipino-Text-Benchmarks-master/sample_sweep.yaml
wandb agent USERNAME/PROJECT_NAME/SWEEP_ID
```

## 7 7. Results and Performance

---

### 7.1 7.1 Model Comparison (qualitative)

- Test Accuracy: ELECTRA ~86% on benchmarks; BiLSTM depends on split and threshold tuning.
- Training Time: BiLSTM ~1–2 hours; ELECTRA ~30–60 mins (hardware dependent).
- Model Size: BiLSTM tens of MB; ELECTRA ~300+ MB.
- GPU Memory: BiLSTM ~2–4 GB; ELECTRA ~8–16 GB.

### 7.2 7.2 Target Achievement

Project Requirement:  $\geq 50\text{--}60\%$  overall accuracy. - ELECTRA: ~86% (Target met with margin). - BiLSTM: Val Acc ~0.843; F1 ~0.885 in a prior run (threshold 0.75 reduced false positives).

### 7.3 7.3 Detailed Performance Analysis - BiLSTM

Error Analysis: - Common False Positives: Sarcasm or strong opinions without hate intent. - Common False Negatives: Subtle/implicit hate; complex code-switching. - Challenging Cases: Sarcasm, irony, context-dependent hate.

### 7.4 7.4 Model Strengths and Limitations

BiLSTM Strengths: - Efficient resource usage; fast inference; captures sequential patterns. BiLSTM Limitations: - May struggle with very long-range dependencies; tuning sensitive. ELECTRA Strengths: - Strong attention to context and superior accuracy.

## 8 8. Tools and Technologies

---

### 8.1 8.1 Deep Learning Framework

- PyTorch for model building, training, and inference.

### 8.2 8.2 Data and Text Processing

- pandas and numpy for data handling and numerical ops.
- For this repository: preprocessing and tokenization are primarily handled in `Filipino-Text-Benchmarks-master/utils/data.py` for transformer models. For the BiLSTM baseline, a prior `text_preprocessing.py` module handled URL/mention handling and tokenization.

## 8.3 Model Training Modules

- `torch.nn` (LSTM, Linear, Dropout); `torch.optim` (Adam/AdamW).

## 8.4 Evaluation and Visualization

- scikit-learn for metrics (accuracy, precision, recall, F1); optional matplotlib for curves.

## 8.5 Development Environment

- Python 3.8+; Git & GitHub for version control.

## 8.6 GUI Applications

- Prior baseline included Tkinter GUI apps (e.g., `gui_app.py`, `social_media_app.py`). These are optional and not part of the ELECTRA training repo; still compatible conceptually for inference UIs.

# 9 Model Deployment

---

## 9.1 Inference Pipeline

The deployment system follows a structured pipeline: 1. Text Input → via GUI, CLI, or service. 2. Preprocessing → clean and tokenize (HF tokenizer for ELECTRA; vocab indices for BiLSTM). 3. Sequence Conversion → integer indices. 4. Padding → to 128 tokens (corrected; not 100). 5. Model Prediction → forward pass. 6. Threshold Application → tuned threshold (e.g., 0.75) for BiLSTM; default argmax for ELECTRA classifier. 7. Output → class label and confidence.

## 9.2 Application Interfaces

- Desktop GUI (baseline): real-time single-text testing.
- Social Media Simulation (baseline): batch moderation demo.

## 9.3 Deployment Files

- ELECTRA fine-tune: checkpoint directory from `train.py` (HF-compatible).
- BiLSTM baseline: `best_bilstm_model.pt`, `vocabulary.pkl`, architecture module.

# 10 Conclusion

---

## 10.1 Project Achievements

- Implemented BiLSTM baseline and ELECTRA fine-tuning path.
- Exceeded minimum accuracy target; ELECTRA ~86% on benchmark; BiLSTM achieved strong validation metrics with threshold tuning.
- Provided working training + evaluation pipeline; optional GUI concepts.
- Conducted dataset validation and bias assessment.

## 10.2 10.2 Key Findings

- BiLSTM effective for sequential classification under resource constraints.
- Dropout (0.3) important for generalization.
- Broader, diverse data improves handling of code-switching.

## 10.3 10.3 Future Improvements

- Explore attention mechanisms, additional Filipino transformers, and better calibration.
- Expand datasets across regions/dialects; consider emoji/sentiment features.

## 10.4 10.4 Lessons Learned

- Preprocessing/tokenization consistency matters, especially for bilingual content.
- Iterative tuning and careful thresholds are key.
- Context-dependence remains challenging; cultural cues are critical.

# 11 11. Practical Notes

---

- HuggingFace auth: `huggingface-cli login` before fetching Hub datasets.
- Boolean CLI args use `true/false` strings in `train.py`.
- Special tokens must match training-time preprocessing (e.g., `[LINK]`).

# 12 12. Reproducibility

---

- Set seeds; pin versions where feasible.
- Keep data splits and cache keys stable between runs.
- Record CLI args and environment (logs or W&B).

Filipino Hate Speech Detection — Revised (2025-12-13)